

# GPU Programming

## Introduction to CUDA Programming

### Learning Objectives

1. Write kernel code for vector addition
2. Explain basic CUDA programming terminologies
3. Explain vector addition CUDA code example

### Cuda Code Example: Vector Add

```
// Kernel code
__global__ void vectorAdd(const float* A, const float* B, float* C, int numElements) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;

    if (i < numElements) {
        C[i] = A[i] * B[i] + 0.0f;
    }
}
```

```
// Host code
int main() {
    // Allocate the device input vector A
    float* d_A = NULL;
    err = cudaMalloc((void**) &d_A, size);

    // Copy the host input vectors A and B in host memory to the device input
    // vectors in device memory
    printf("Copy input data from the host memory to the CUDA device\n");
    err = cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);

    vectorAdd<<<blocksPerGrid, threadsPerBlock>>>>(d_A, d_B, d_C, numElements);

    // Copy the device result vector in device memory to the host result vector
    // in host memory
    printf("Copy output data from the CUDA device to the host memory\n");
    err = cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
}
```

1. Host code is executed on CPUs
2. Kernel code is invoked with «<...>»
3. Kernel code is executed on GPUs

### SPMD

1. GPU kernel code is single program multiple data
  - All threads execute the same program
  - There is no execution order among threads
  - But we need to make each thread execute different data

### threadIdx.x

1. Even though each thread executes the same program, each thread has a unique identifier (each thread has built-in variable that represents the x-axis coordinate)
  - threadIdx.x

## Vector Add with Thread IDs

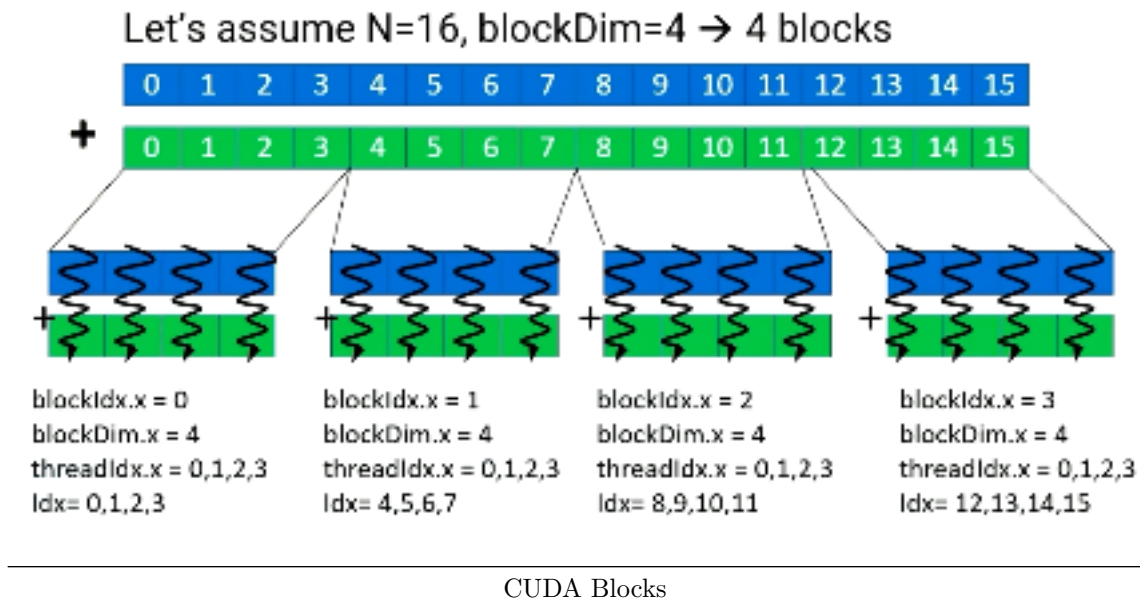
1. Each thread operates on a different element of vector data
  - `int idx = threadIdx.x;`
  - `c[idx] = a[idx] + b[idx];`

## Execution Hierarchy

1. A group of threads forms a block
  - Group of threads is also called a warp
    - Warp is a microarchitecture concern
      - \* Pure hardware decision, programmer doesn't need to be aware
  - Block is a critical component of CUDA
2. CUDA block: A group of threads that are executed concurrently
3. Data is divided by block
4. For now, let's just assume that each block is executed on each GPU SM
  - SM: Streaming Multiprocessor
5. No ordering among CUDA block execution

## Example of Data Indexing

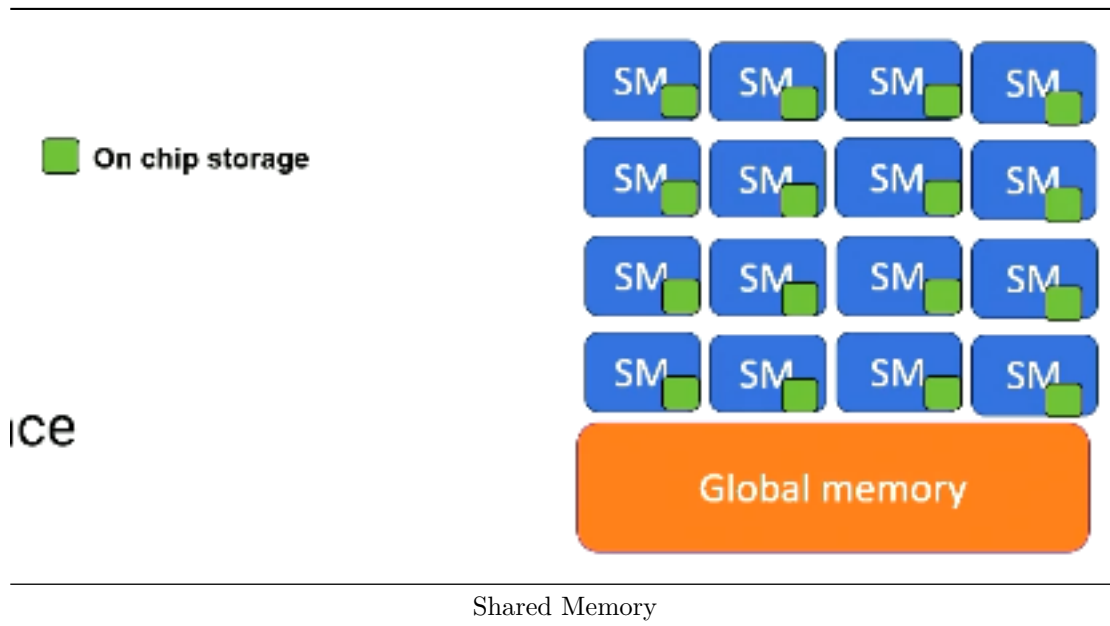
1. Thread Index
  - `threadIdx.x`, `threadIdx.y`, `threadIdx.z`
2. Block Index
  - `blockIdx.x`, `blockIdx.y`
3. Block Dimensions
  - `blockIdx.x * blockDim.x + threadIdx.x`



## Shared Memory

1. Scratchpad memory
2. Software-controlled memory space
3. Use   shared
4. On chip storage  $\rightarrow$  faster access compared to global memory

5. Accessible only within a CUDA block (later GPUs allow different policy)



### Execution Ordering: Threads and Blocks

1. No predefined ordering among threads or CUDA blocks

### Thread Synchronizations

1. Wait for all threads to complete tasks
  - Typical usage: ‘\_\_syncthreads()’ within a CUDA block
  - Interblock synchronization: achieved through different kernel launches

### Typical Usage with Thread Synchronizations

1. Typical computation flow
  - Load -> Compute -> Store
  - Load to the shared memory from global memory
  - Compute with the data in the shared memory
  - Store the results into global memory
2. BSP programming model (bulk synchronous parallel)

### Kernel Launch

1. Launch kernel from the host code using ‘<<<>>>’ syntax
  - kernelName<<<numBlocks, threadsPerBlock>>>(arguments)
2. Total number of threads in a kernel = numBlocks \* threadsPerBlock
3. Grid consists of multiple blocks, block consists of multiple threads
4. Multiple different tasks -> multiple kernels
5. Each kernel (sortKernel, addKernel, storeKernel) handles different tasks sequentially

```
int main() {
    sortKernel<<<1,1>>>(d_data, dataSize);
    addKernel<<<gridDim,blockDim>>>(d_data, d_data, d_result, dataSize);
    storeKernel<<<1,1>>>(d_result, dataSize);
}
```

## Memory Space and Block/Thread

	Scope	Note
Shared Memory	Within one CUDA block	
Global Memory	All blocks in a kernel	
Local Memory	Within a CUDA thread	
Constant Memory	All blocks in a kernel	Read only (3D-graphics)
Texture Memory	All blocks in a kernel	Read only (3D-graphics)

1. Information sharing is limited across “CUDA execution” hierarchy:
  - Data in the shared memory is visible only within one CUDA block
    - Data in shared memory can stay only in one SM
    - Data in the shared memory of one CUDA block needs explicit communication (later CUDA supports thread block cluster to allow sharing)

## Summary

1. Learned execution hierarchy, shared memory, global memory, kernel launch, thread synchronization

## Occupancy

### Learning Objectives

1. Explain the concept of occupancy
2. Determine how many CUDA blocks can run on one Streaming Multiprocessor

### How Many CUDA Blocks on One Stream Multiprocessor?

1. So far, we assumed on CUDA block per SM
  - In reality, there are multiple blocks per SM

## Occupancy

1. How many CUDA blocks can be executed on one SM is decided by the following parameters:
  - # of registers
  - shared memory
  - # of threads
2. Exact hardware configurations are varied by GPU microarchitecture
3. Example
  - Each SM can execute 256 threads, 64K registers, 32 KB shared memory
    - One CUDA block: 32 threads and 2KB shared memory and each thread has 64 registers
    - Constrain by running threads:  $256/32 = 8$  CUDA blocks
    - Constrain by register size:  $64 * 1024 / (64 * 32) = 32$  CUDA blocks
    - Constrain by shared memory =  $32KB / 2KB = 16$  CUDA blocks
  - Final answer: 8 CUDA blocks per SM

## Number of Threads Per CUDA Block

1. Host sets the number of threads per CUDA blocks
2. Set up at the kernel launch time
3. # of registers per CUDA block
  - Compiler sets how many architecture registers are needed for each CUDA block (will be covered in later part of the course)
  - Typical CPU ISAs, # of registers per thread is fixed (e.g., 32), but in CUDA, this is flexible
4. Shared memory size is determined from the code

- `__shared__ int sharedMemory[1000];` -> 4000B

### Why Do We Care about Occupancy?

1. Higher occupancy means more parallelism
2. Utilization: Utilizing more hardware resources generally leads to better performance
3. Exceptional cases will be studied in later lectures

### Summary

1. Occupancy depends on # of threads per CUDA block, register, and shared memory usages
2. Higher occupancy generally leads to more parallelism


## Host Code

### Learning Objectives

1. Write the host code for vector addition operation
2. Explain the management of global memory

### Managing Global Memory

1. Shared memory: Visible within CUDA block
2. Global memory: Visible among all CUDA blocks
3. A host program needs explicit APIs to manage device memory (global memory)



### CUDA API

1. `cudaMalloc`: allocate memory on the GPU
2. `cudaMemcpy`: transfer data between CPU and GPU
3. Unified memory: Eliminates the need for explicit data copies
  - Covered in later lectures

### Summary

1. Host launches kernel and manages device memory including allocation and data transfer


## Stencil Operation with CUDA

### Learning Objectives

1. Be able to write a stencil code with CUDA shared memory and synchronization

### Recap: Stencil Operations

1. A common pattern in HPC, perform computation with neighboring data
2. Each CUDA block performs different calculation
  - Each CUDA block operates each color zone
3. Operations are fully parallel
4. Data decomposition: Each thread handles one element of the stencil operation
5. One element will be used at least 4 times
  - Good candidates for caching



---

Stencil Operations

---

### Stencil Operations: Operations per Thread

1. Stage 1: Load data from global memory to shared memory (on-chip storage)
2. Stage 2: Perform stencil operations
3. Stage 3: Write results back
4. Thread synchronization between stage 1 and stage 2

### How about Boundaries?

1. C position cannot access elements 1 and 3
2. Solution: bring neighboring elements together and pad with zeros outside of the original size

```
// load data into shared memory
__shared__ float sharedInput[sharedDim][sharedDim];
```

```
int sharedX = threadIdx.x + filterSize / 2;
int sharedY = threadIdx.y + filterSize / 2;
```

### Use if-else to Check Boundary or Not

1. Load different values on boundaries

```
if (x >= 0 && x < width && y >= 0 && y < height) {
    sharedInput[sharedY][sharedX] = input[y * width + x];
} else {
    sharedInput[sharedY][sharedX] = 0.0f; // handle boundary conditions
}
```

2. Perform computations only on inner elements

```
// apply the filter to the pixel and its neighbors using shared memory
for (int i = 0; i < filterSize; i++) {
    for (int j = 0; j < filterSize; j++) {
        result += sharedInput[threadIdx.y + i][threadIdx.x + j] * filter[i][j];
    }
}
```

### Summary

1. Explored writing a stencil operation kernel with shared memory
2. Discussed how to handle boundary elements with if-else statements